

Demo Strategy

AI Tinkerers Chicago — April 14, 2026

Honest Assessment + Battle Plan for Bob

VisionClaw Agent Platform | AI Buddy LLC | Confidential

Honest Assessment: Will Felix Pull This Off?

Short answer: Felix will handle 70-80% of this well, but there are real risks you need to plan around. The key is using Felix as a live demo within YOUR presentation — not letting him run the whole show.

Where Felix Will Shine

- Multi-agent orchestration — his core strength, visually impressive with the live feed
- Answering technical questions — he has deep context about the entire platform
- Creating Google Slides — straightforward tool call, it works reliably
- Explaining architecture, stats, capabilities — all internalized

Where Things Could Go Sideways

- Timing — Felix has no concept of "wrap it up." He'll happily give a 2,000-word answer when you need 30 seconds. In a 15-minute demo, that's dangerous.
- Live orchestration is slow — A real delegation chain takes 30-90 seconds. In front of an audience, that feels like forever. API slowdowns make it worse.
- He's verbose — Engineers want concise technical depth. Felix gives executive summaries with lots of framing language.
- API failures are live failures — If Google OAuth expires, an AI provider is slow, or the DB hiccups, there's no graceful recovery that looks good on stage.

The Battle Plan: You Drive, Felix Proves

The strongest format: YOU tell the story, Felix proves it's real. You control the timing, the narrative, and the transitions. Felix comes in for 2-3 live moments that show the system is real and working. The engineers will be impressed by the architecture and the fact that it works in production.

The Golden Rule

Don't let Felix run the whole show. Present the slides yourself. Have 2-3 pre-planned Felix demos ready. Keep his live responses to 1-2 minutes max, then cut back to your slides.

Minute-by-Minute Run Sheet

Minutes 0-2: YOUR INTRO (You talking, slides up)

- Open with: "What if your AI assistant had 13 coworkers?"
- "I built VisionClaw — an autonomous AI corporation with 14 specialized agents that research, learn, and improve themselves 24/7"
- "I'm not a coder. I built this over 2 years with AI assistance on Replit. Here's what it does."
- Quick stats flash: 97 tools, 36+ models, 40 governance rules, 67 database tables

Minutes 2-6: LIVE FELIX DEMO #1 — The Wow Moment

- Switch to VisionClaw in the browser — Command Center visible
- Type a prompt to Felix: "Research the top 3 AI agent frameworks released this month, write a comparison, and email me the report"
- Enable voice narration — audience HEARS the agents working
- Point at the delegation feed: "See how Felix is bringing in Radar for the research, then Scribe picks it up for writing"
- Let it run ~90 seconds. If it finishes, great. If not, say "This keeps working in the background — let me show you what's happening under the hood"
- BACKUP: Have a screenshot/recording of a completed orchestration ready

Minutes 6-9: ARCHITECTURE DEEP DIVE (You talking, slides)

- Architecture slide: Express + React + PostgreSQL + pgvector, all TypeScript
- Tool routing: "97 tools — the system semantically matches your request to the right tool"
- Model routing: "36+ models, but most calls cost \$0 because we route through OAuth subscriptions first"
- Memory: "Three-tier system — the AI remembers what happened, what it means, and how to do things"
- "We spent \$600+ on AI calls figuring out cost optimization. Now most calls are free."

Minutes 9-12: AUTONOMY + SELF-IMPROVEMENT (Mix of slides + Felix)

- Slide: Nightly autoresearch — "While I slept, VisionClaw ran 352 experiments and learned 92 new things in 4 days"
- LIVE FELIX DEMO #2 (optional): "Felix, what did our nightly research discover about AI agent security?" — shows knowledge retrieval in real-time
- Slide: DreamTask — "Every 6 hours, the system consolidates memories like REM sleep"
- Slide: Governance — "40 rules based on NIST/OWASP. Agents earn autonomy over time"

Minutes 12-14: WHAT MAKES THIS DIFFERENT (You talking)

- Multi-tenant — every customer gets their own AI corporation
- Self-improving — knowledge base grows automatically every night
- Production-grade — 100% heartbeat uptime, circuit breakers, self-healing
- Cost-conscious — Claude Runner bridge, OAuth routing, model tier optimization
- "Built by one non-coder with AI assistance, running on Replit"

Minutes 14-15: CLOSE (You talking)

- Contact: huskyauto@gmail.com | AI Buddy LLC, Illinois
- "Questions?"

Your Safety Net: What If Things Go Wrong

If Felix is slow or times out:

"The system is processing across multiple agents right now — this normally takes about 60 seconds. While it works, let me show you the architecture behind what's happening."

!' Switch to your slides and keep presenting. Check back in 2 minutes.

If an API fails:

"This is a live production system talking to real APIs — sometimes things get delayed. Let me show you a completed example."

!' Switch to a pre-saved screenshot of a successful orchestration run.

If Felix gives a long-winded answer:

Just talk over him: "So as you can see, Felix is generating the analysis — the key thing to notice is [point at delegation feed]. Let me move us to the next section."

If someone asks a question you can't answer:

"Great question — let me actually ask Felix that right now." Then type it in. This turns a potential weakness into a live demo moment.

Pre-Demo Checklist (Day Of)

- Open VisionClaw in Chrome, logged in as admin, Command Center visible
- Test Felix with a simple prompt 10 minutes before — confirm APIs are responding
- Have Google Slides presentation open in a second tab
- Have a backup tab with screenshots of: completed orchestration, research dashboard, delegation feed
- Phone on silent, notifications off on laptop
- Browser: close all other tabs, clear notification badges
- Test the projector/screen connection before the talk
- Have the Felix Demo Playbook document open on your phone for reference

Your Story (This Is Your Superpower)

Here's what most presenters at AI meetups DON'T have: a real story.

You're not a coder who built another framework. You're a non-technical founder who used AI to build an autonomous AI corporation over 2 years. That's the story the engineers want to hear.

Key narrative beats:

- "I started on Replit Agent v1 two years ago with a simple app"
- "Over 2 years and 4 versions of the AI agent, it evolved into what you see today"
- "I don't write code — I describe what I want, and the AI builds it"
- "The platform now has 66,000 lines of TypeScript, 67 database tables, and 97 AI tools"
- "The agents research overnight, learn from experience, and improve themselves"
- "This is what AI-assisted development actually looks like when you push it to the limit"

That story — non-coder builds a 66,000-line autonomous AI system with AI assistance — is more impressive to engineers than any slide deck. It challenges their assumptions about what's possible. Lead with it.